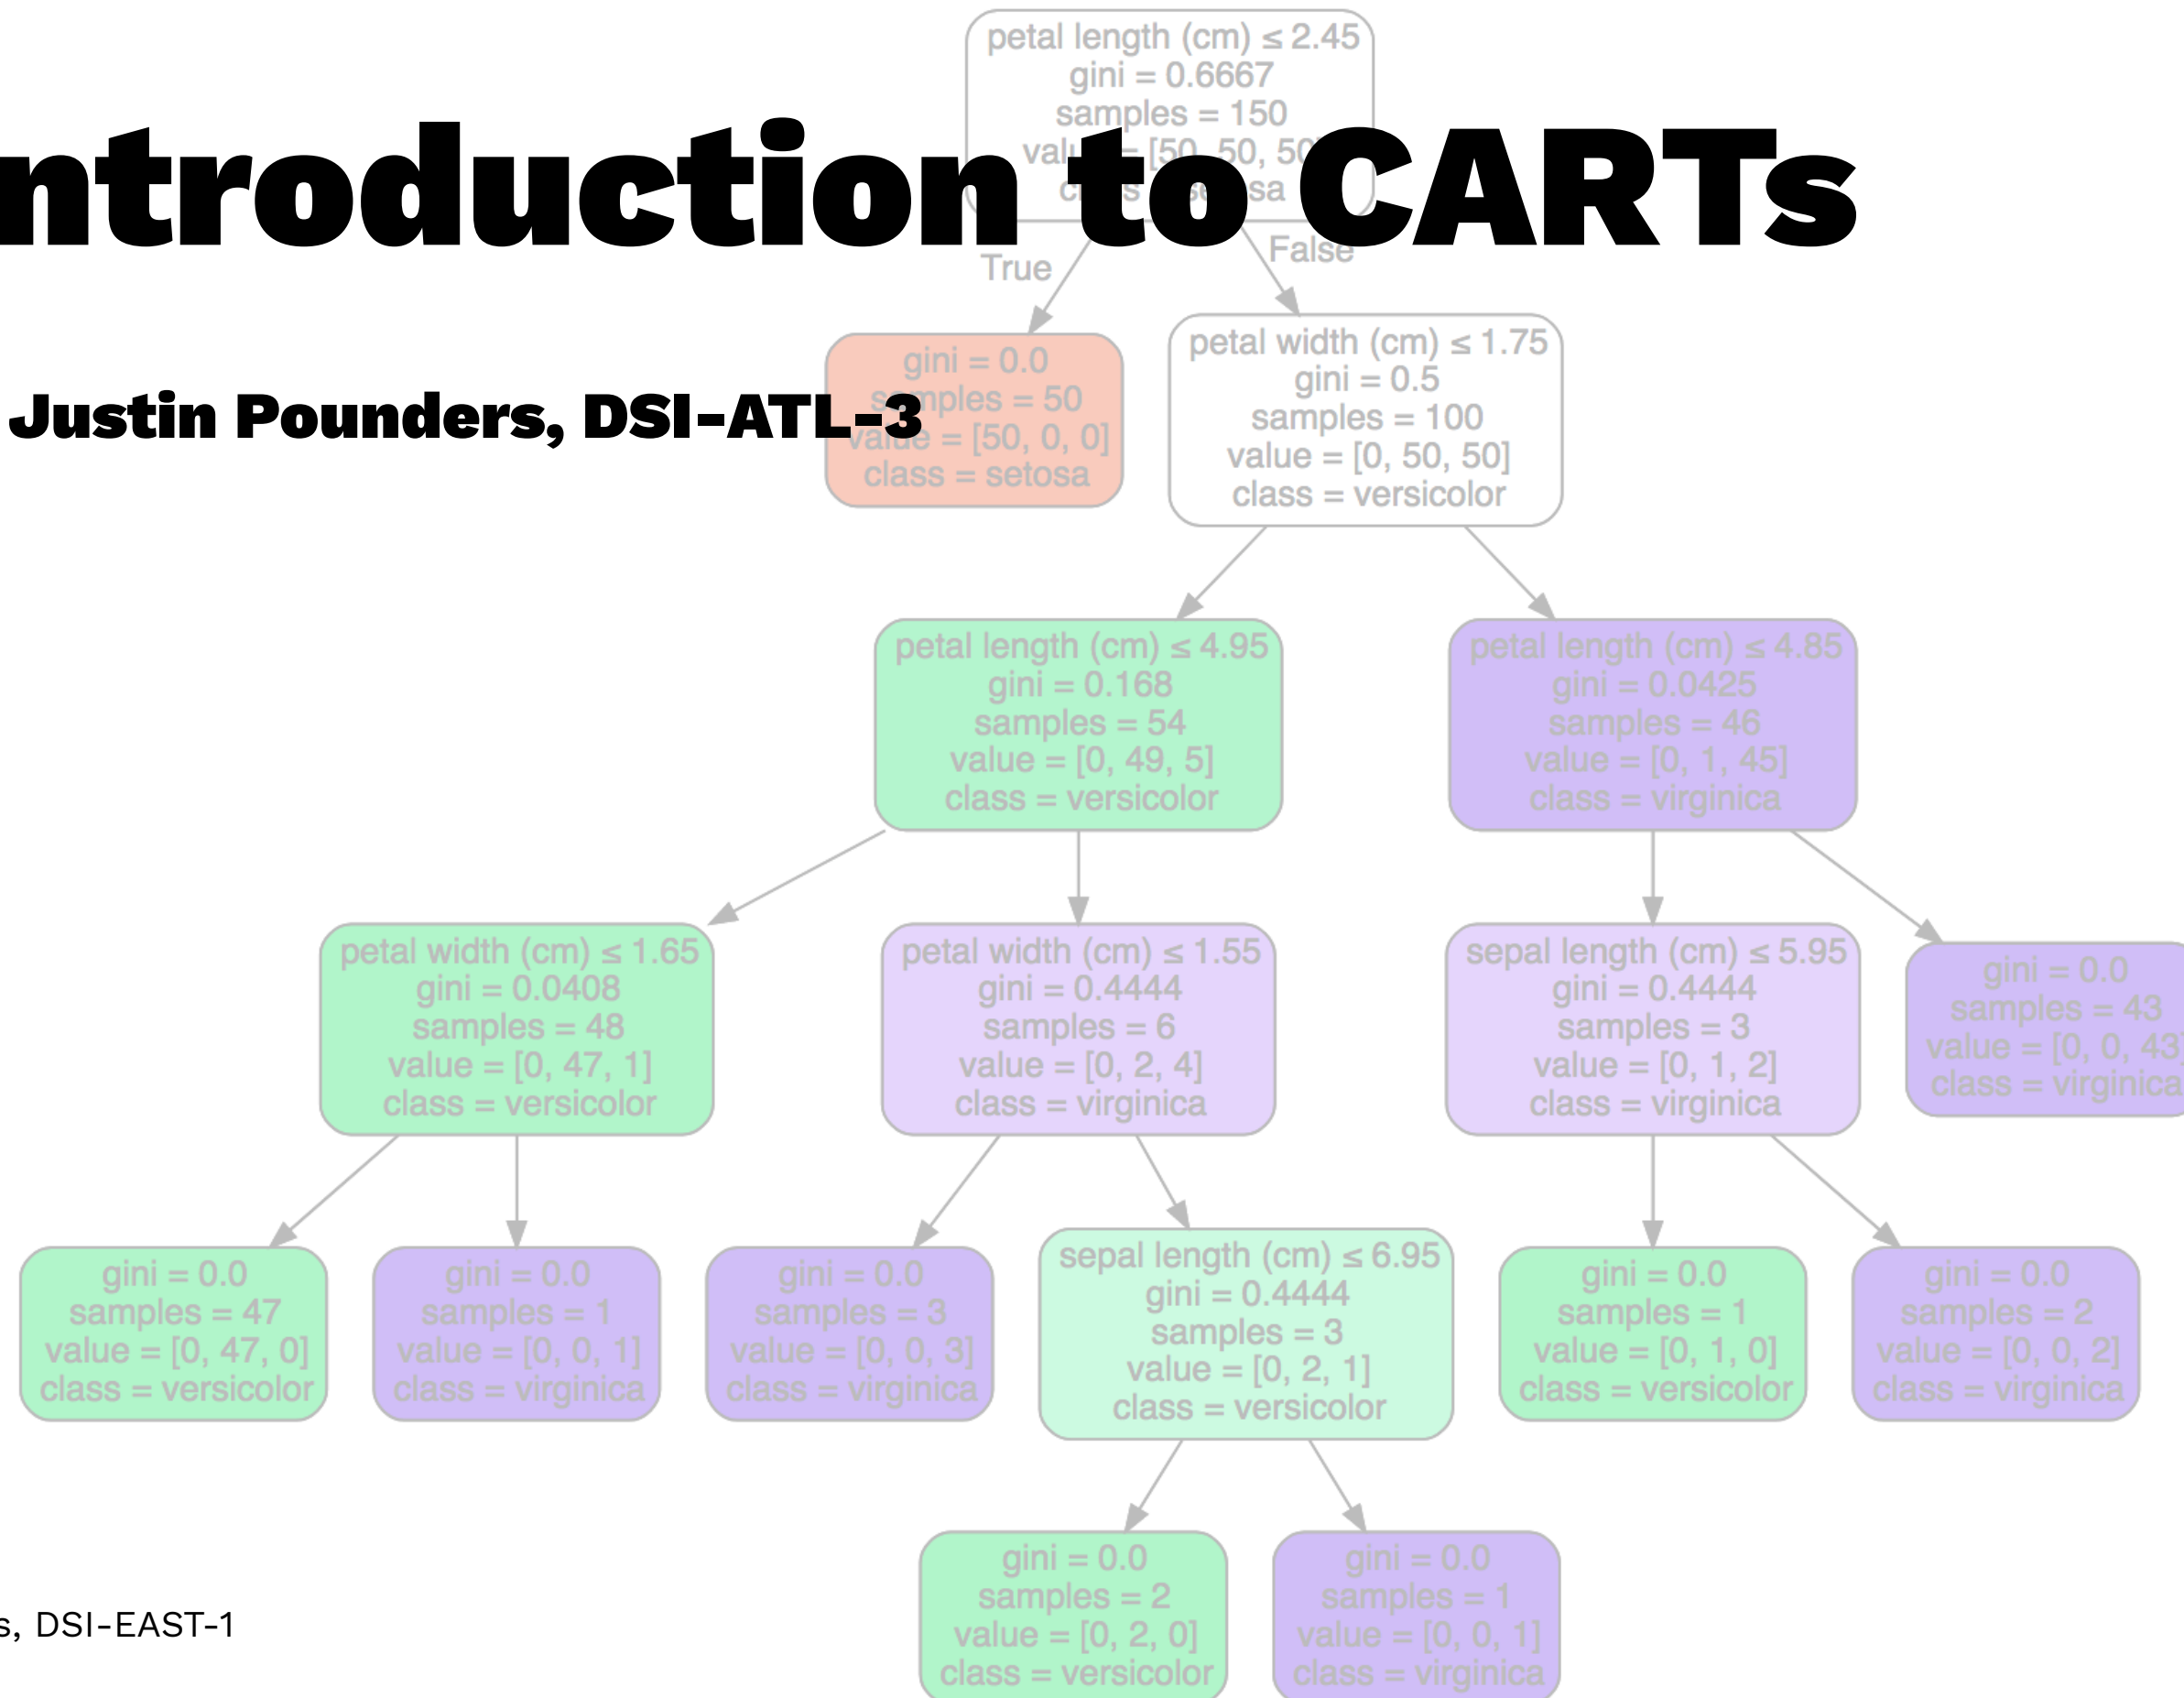




# Introduction to CARTs

**Author: Justin Ponders, DSI-ATL-3**



# Objectives

- Describe how decision tree models work
- Define the concept of purity and information gain
- Describe how model ensembles lead to better performance
- Define boosting and model aggregation (aka bagging)
- Build decision tree and random forest models in sklearn
- Add ensemble bagging to any classification model in sklearn

# **Guess what animal I am thinking of.**

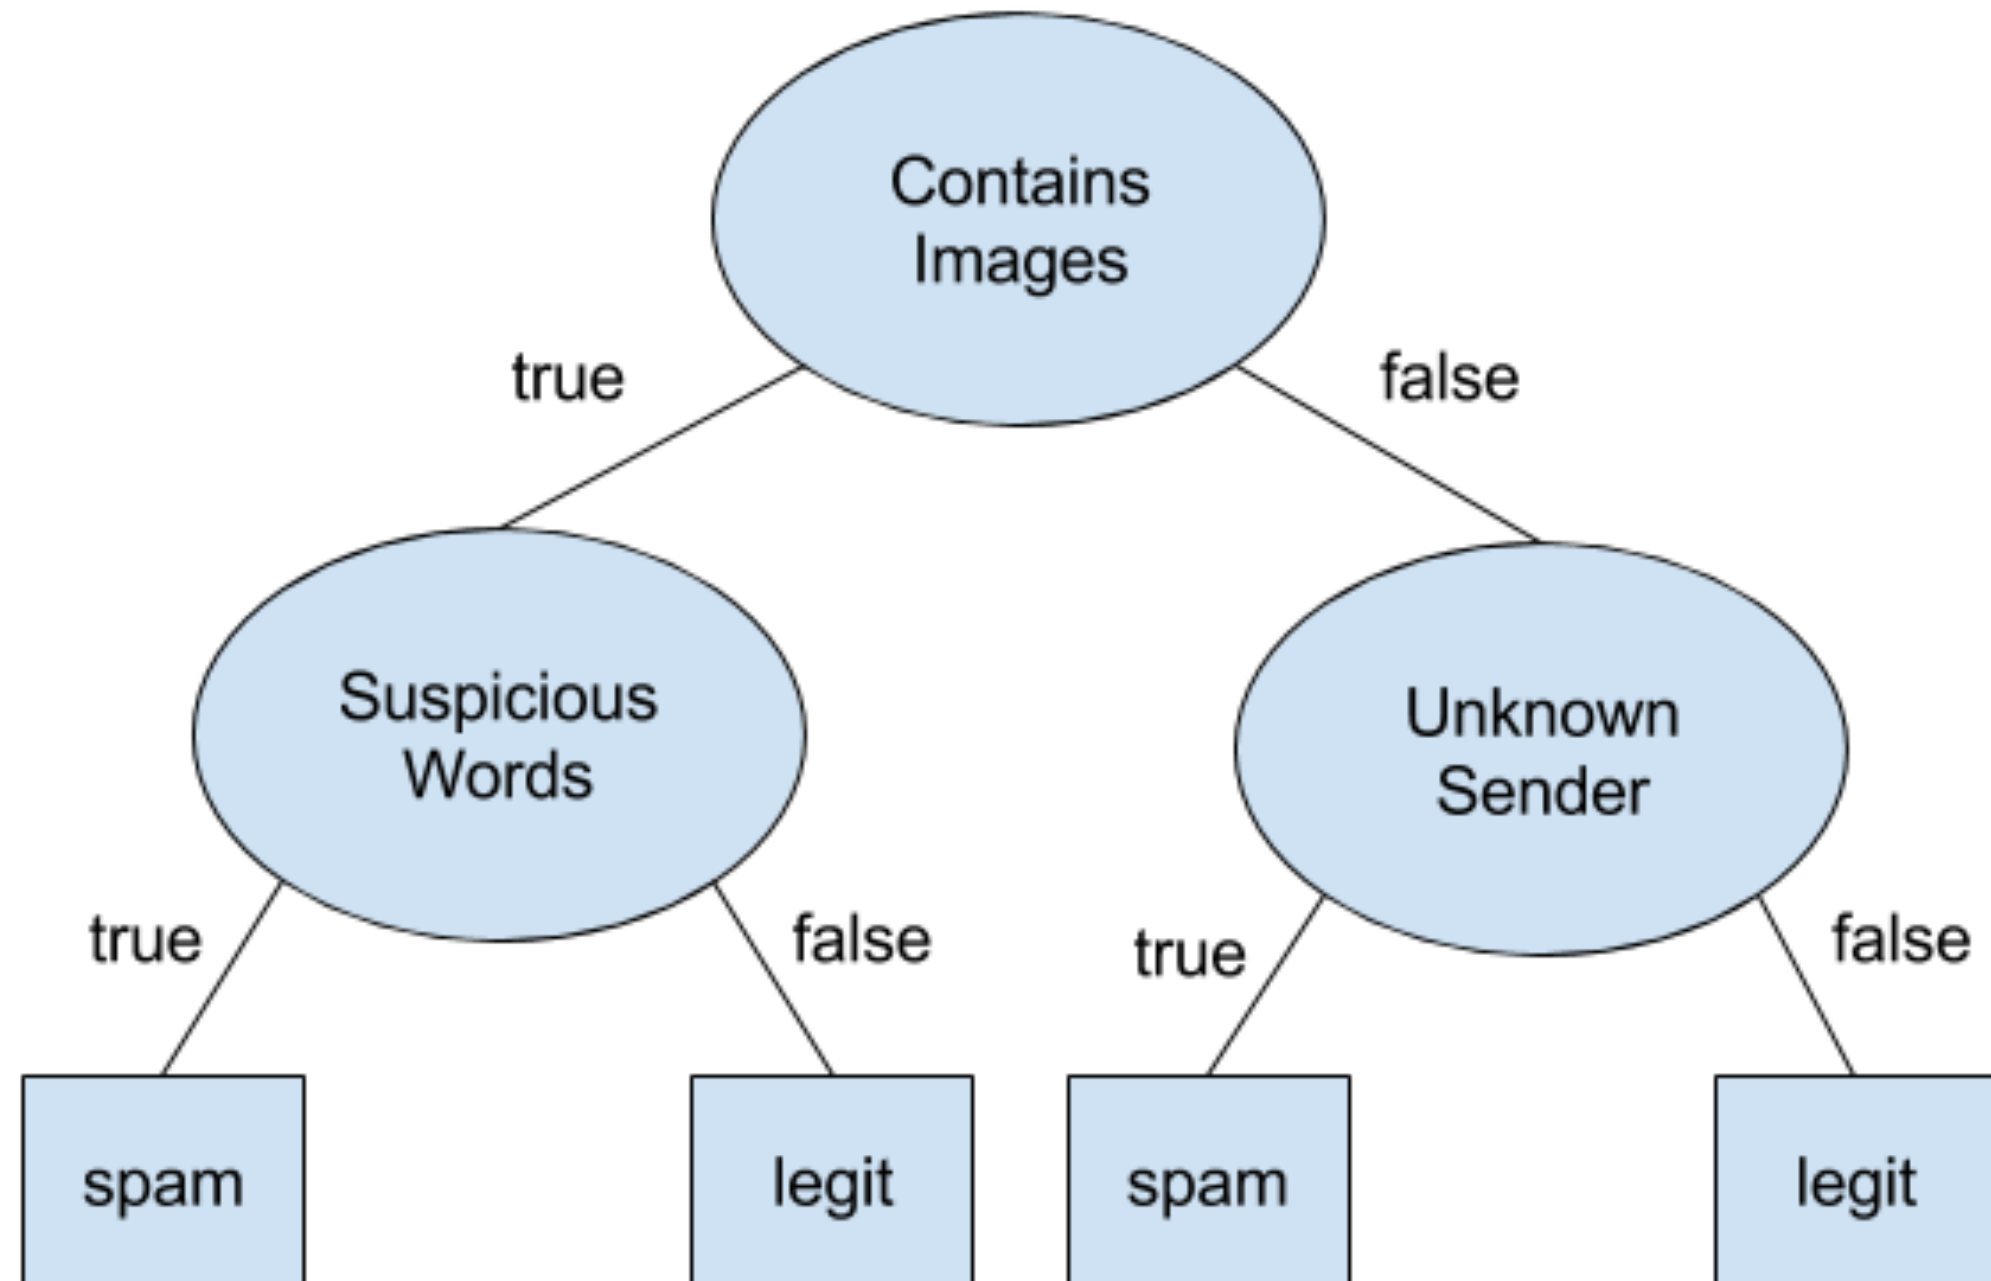
# Intro to CARTs

**CART = Classification and Regression Tree**

**CART = Decision Tree**

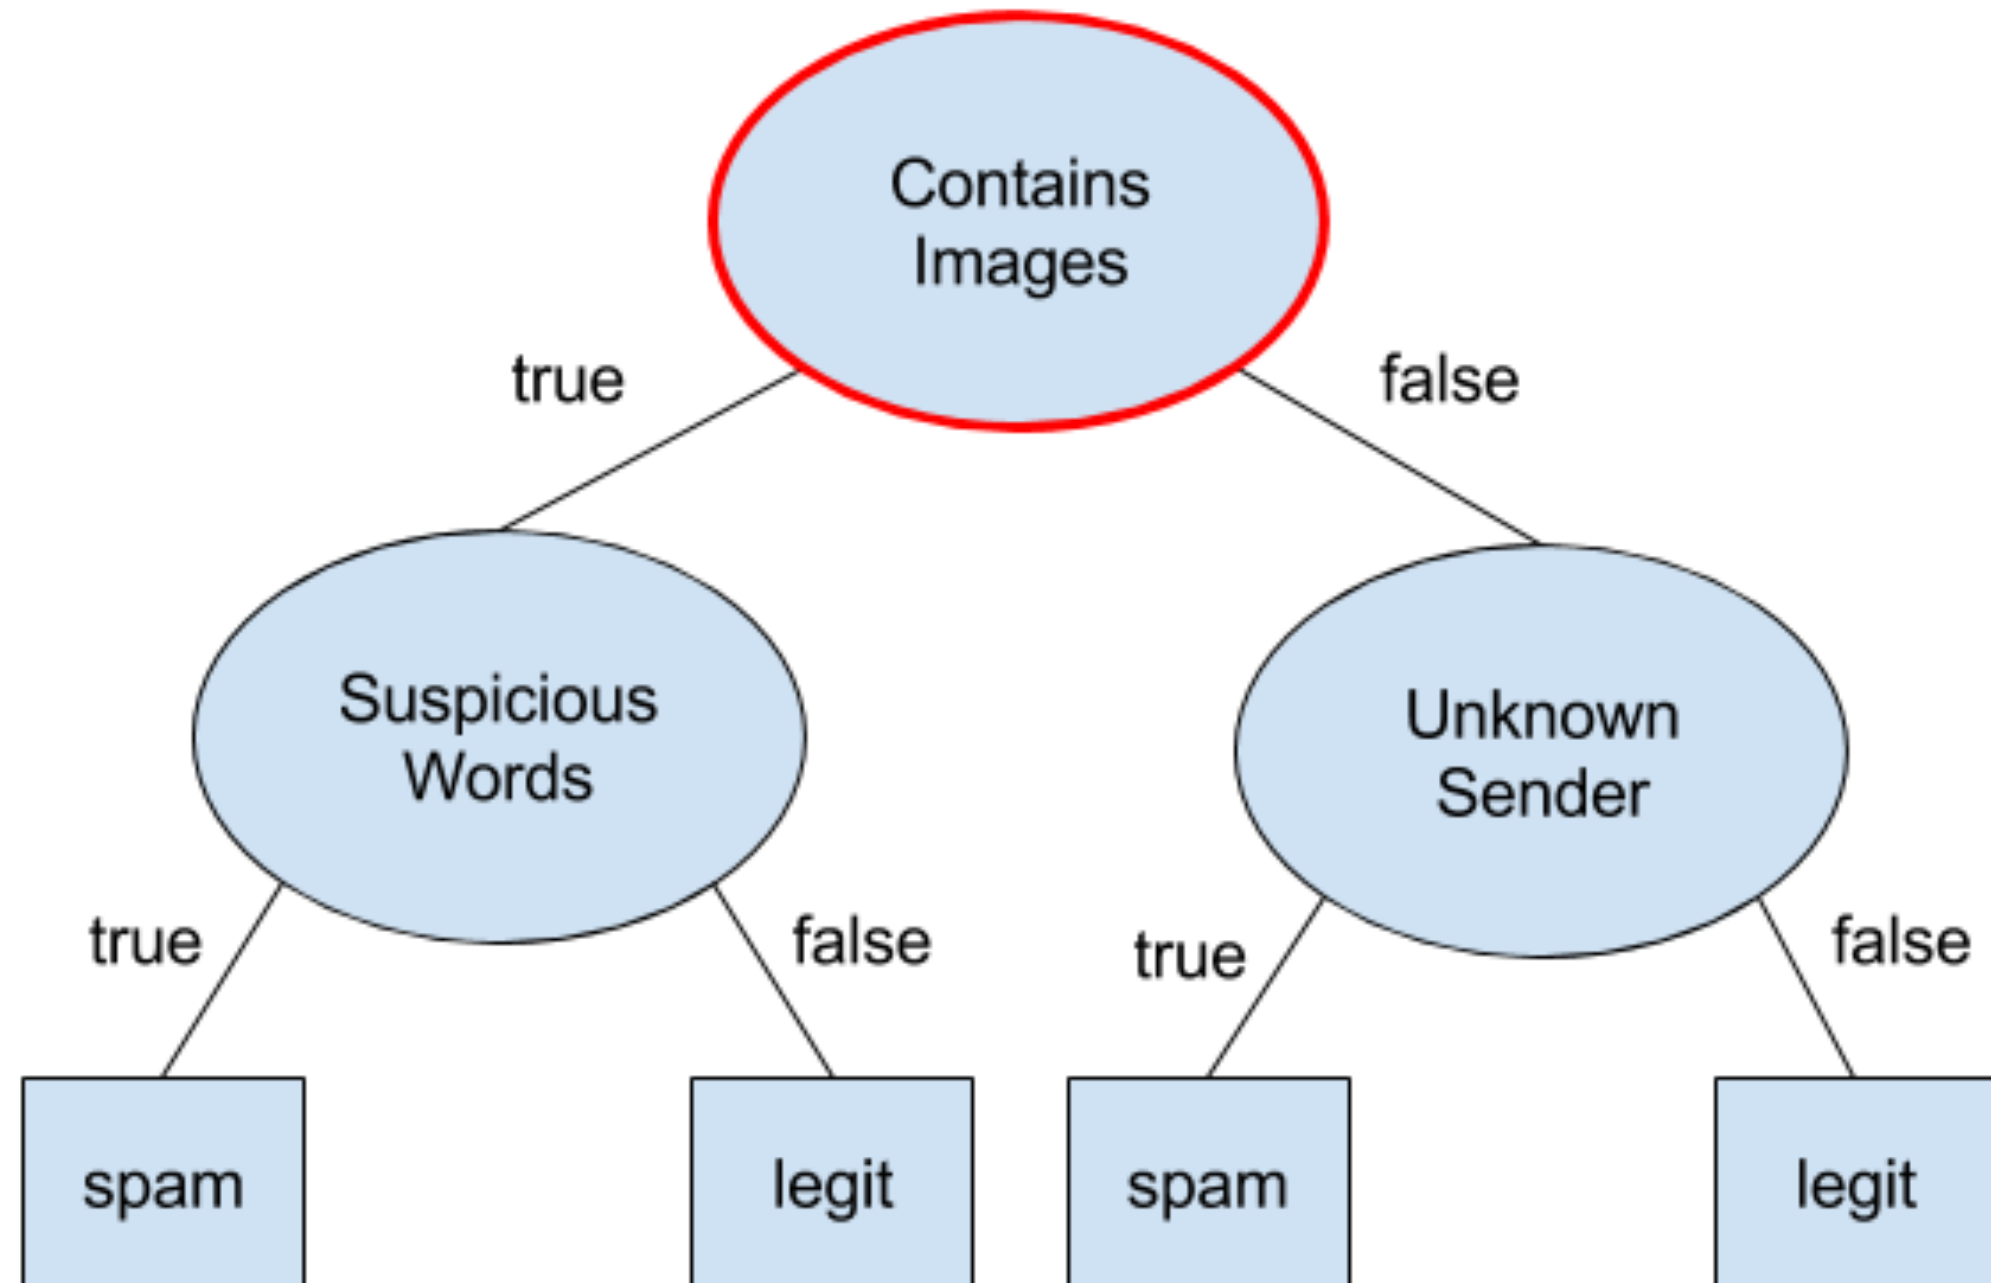
<i>ID</i>	<i>Susp. words</i>	<i>Unknown sender</i>	<i>Contains images</i>	<i>Classifi- cation</i>
1	true	false	true	???

## Example Decision Tree for a Spam Filter



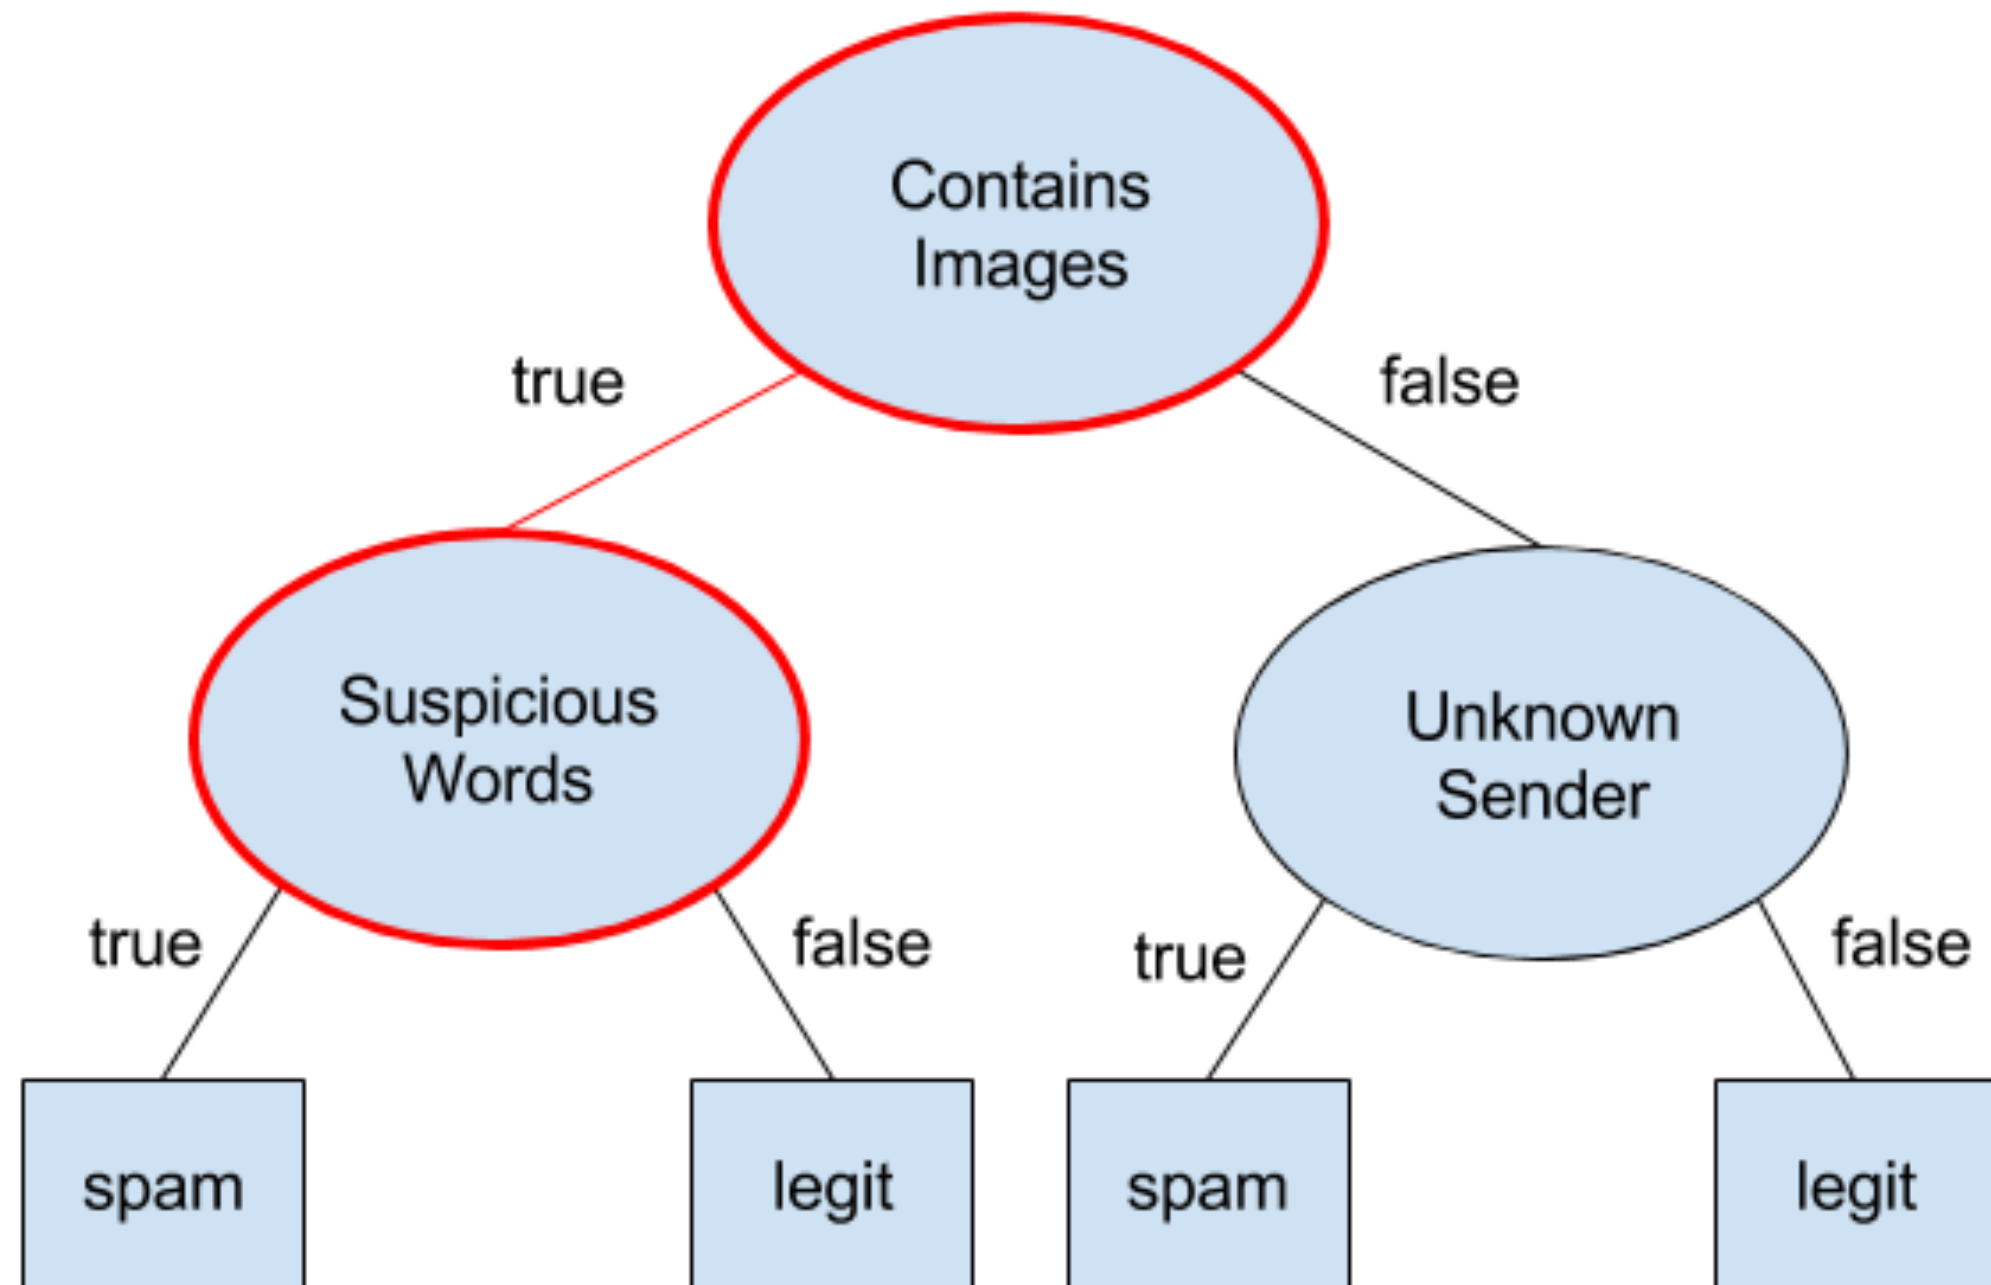
<i>ID</i>	<i>Susp. words</i>	<i>Unknown sender</i>	<i>Contains images</i>	<i>Classification</i>
1	true	false	true	???

## Example Decision Tree for a Spam Filter



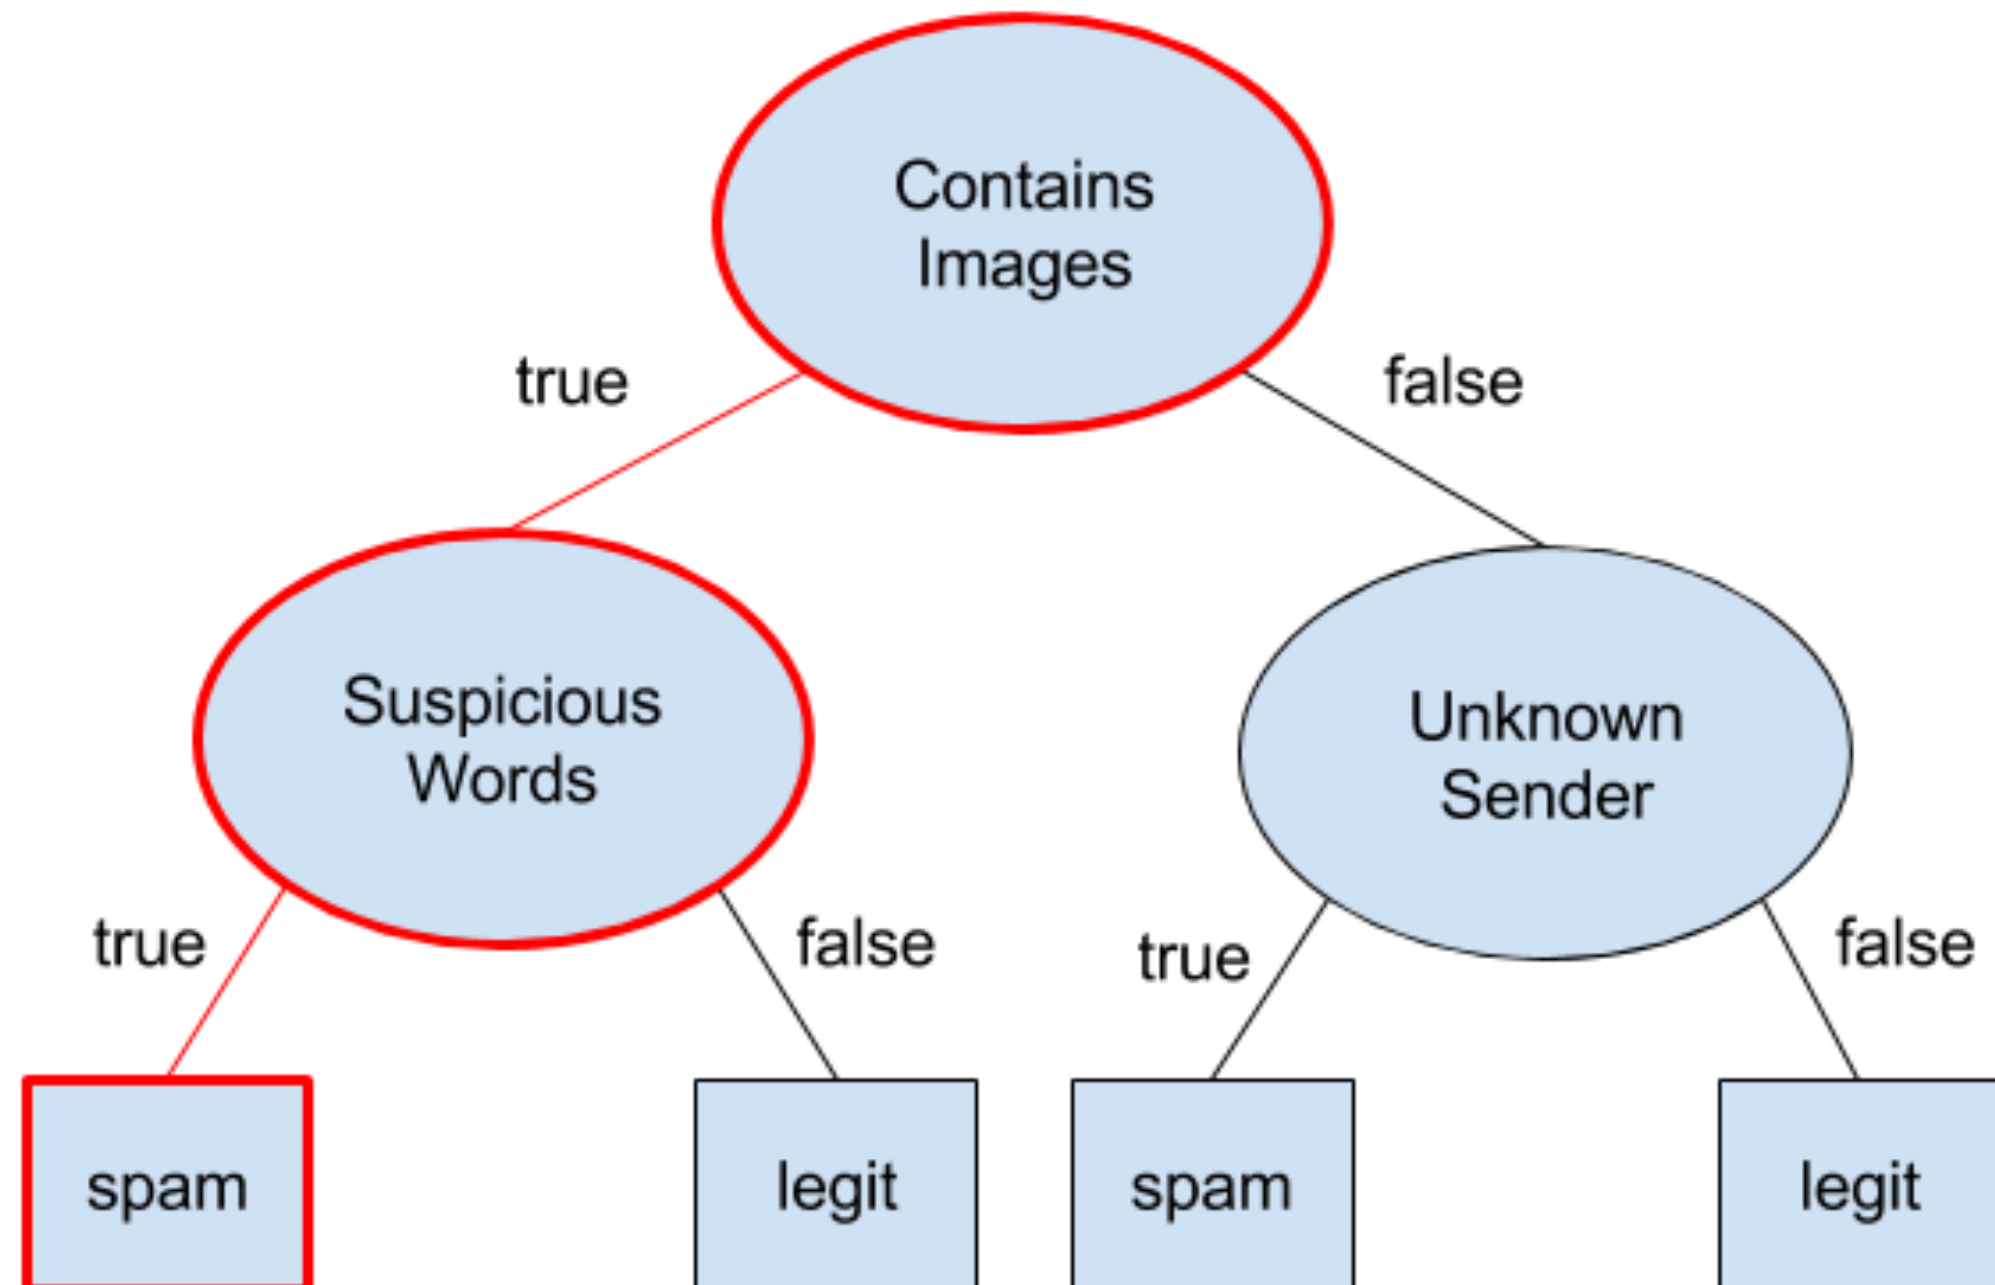
<i>ID</i>	<i>Susp. words</i>	<i>Unknown sender</i>	<i>Contains images</i>	<i>Classification</i>
1	true	false	true	???

## Example Decision Tree for a Spam Filter



<i>ID</i>	<i>Susp. words</i>	<i>Unknown sender</i>	<i>Contains images</i>	<i>Classifi- cation</i>
1	true	false	true	spam

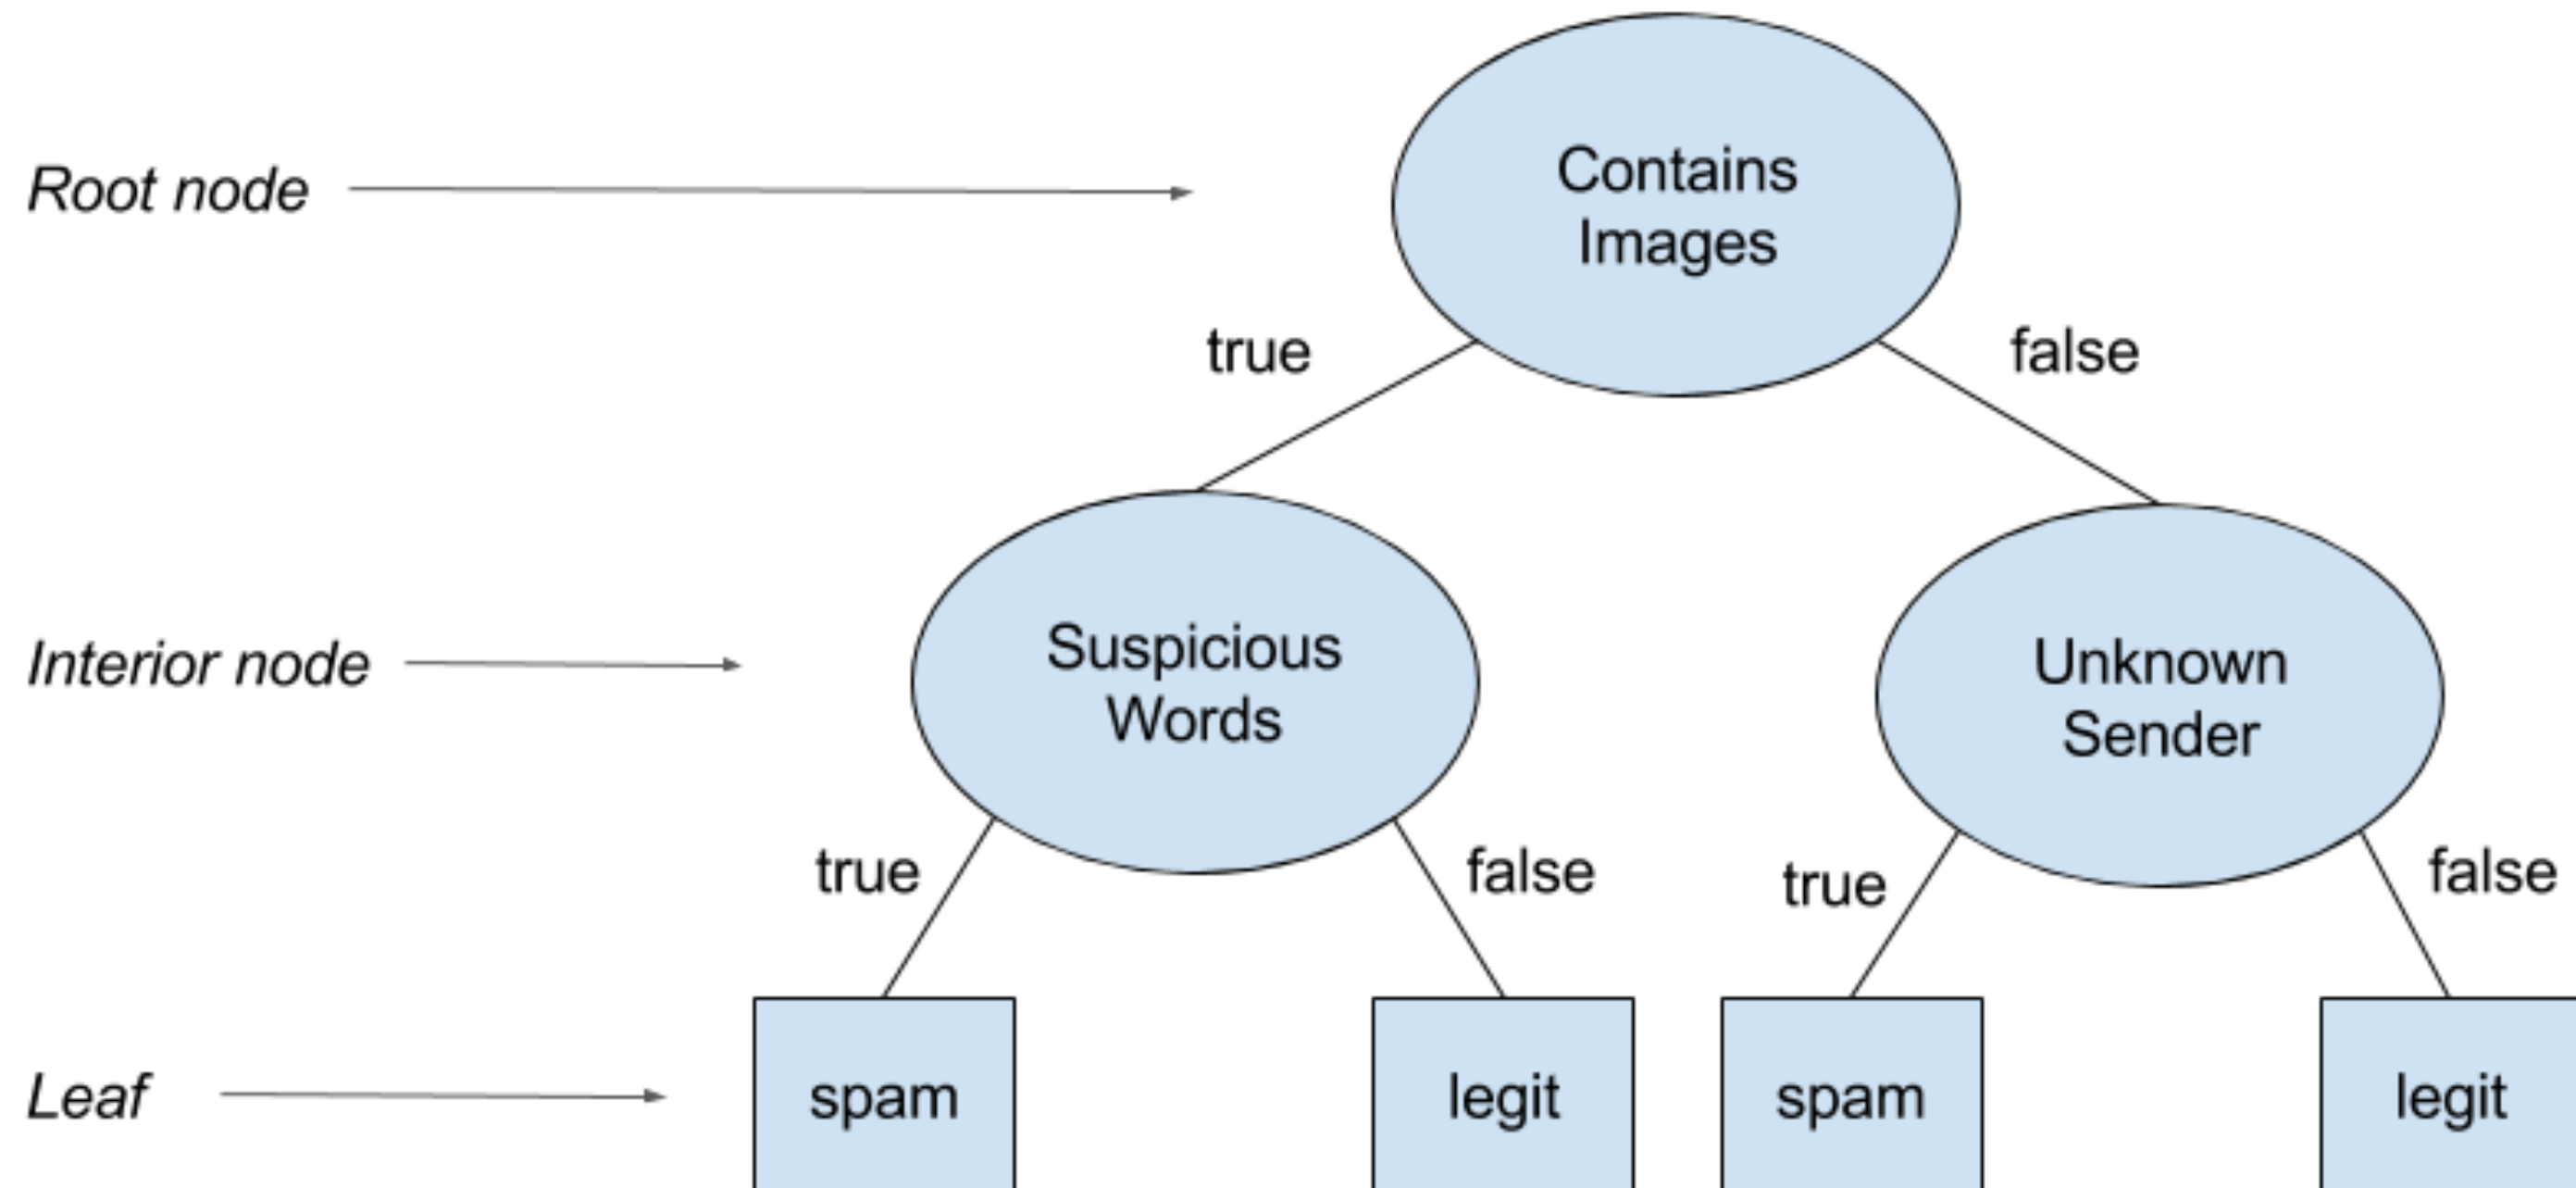
## Example Decision Tree for a Spam Filter





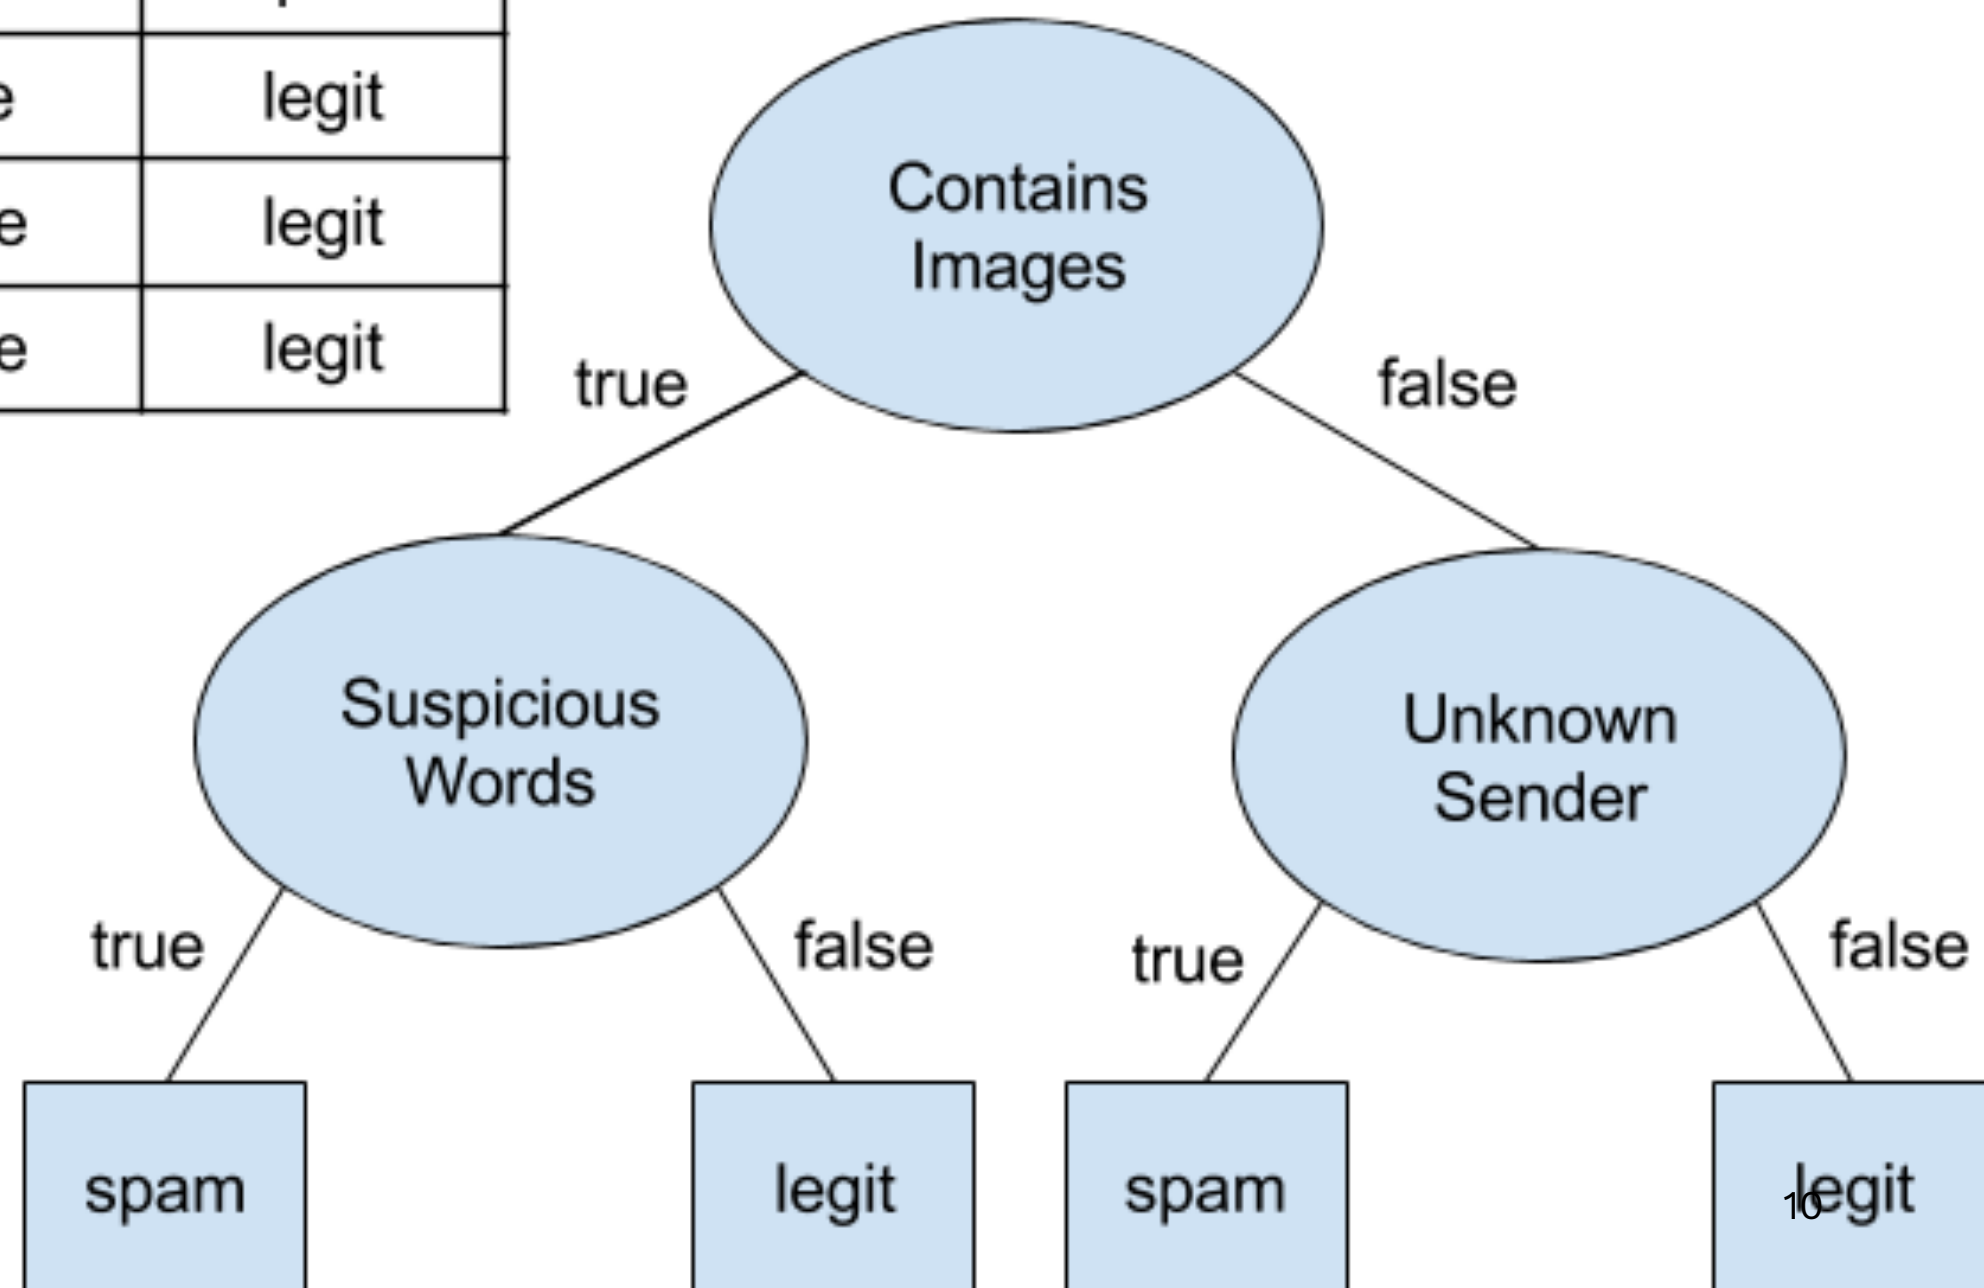
<i>ID</i>	<i>Susp. words</i>	<i>Unknown sender</i>	<i>Contains images</i>	<i>Classification</i>
1	true	false	true	spam

## Example Decision Tree for a Spam Filter




<i>ID</i>	<i>Susp. words</i>	<i>Unknown sender</i>	<i>Contains images</i>	<i>Classifi- cation</i>
1	true	false	true	spam
2	true	true	false	spam
3	true	true	true	spam
4	false	true	true	legit
5	false	false	false	legit
6	false	false	false	legit

How do you build a decision tree?



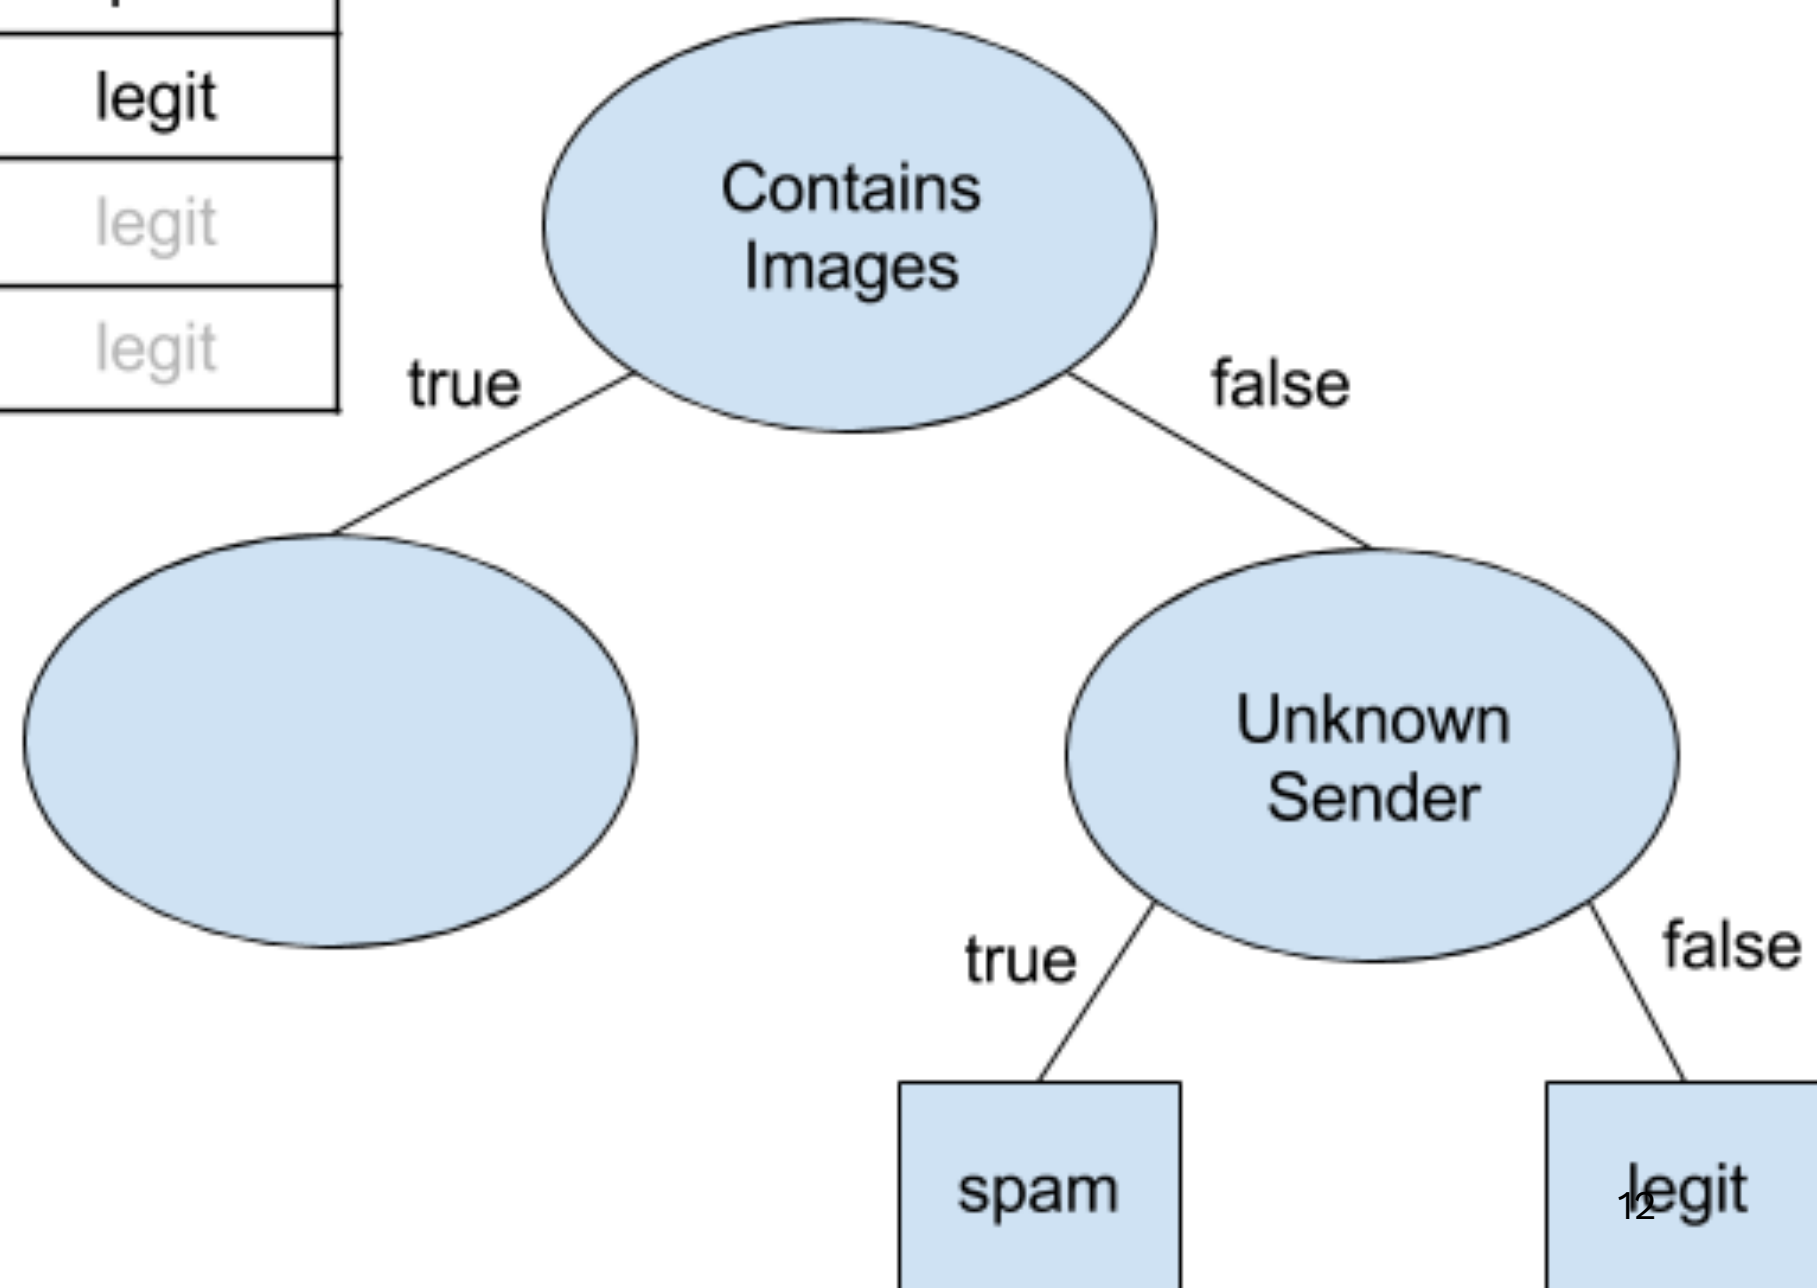
ID	Susp. words	Unknown sender	Contains images	Classification
1	true	false	true	spam
2	true	true	false	spam
3	true	true	true	spam
4	false	true	true	legit
5	false	false	false	legit
6	false	false	false	legit

Each new node generates a subset of the data.



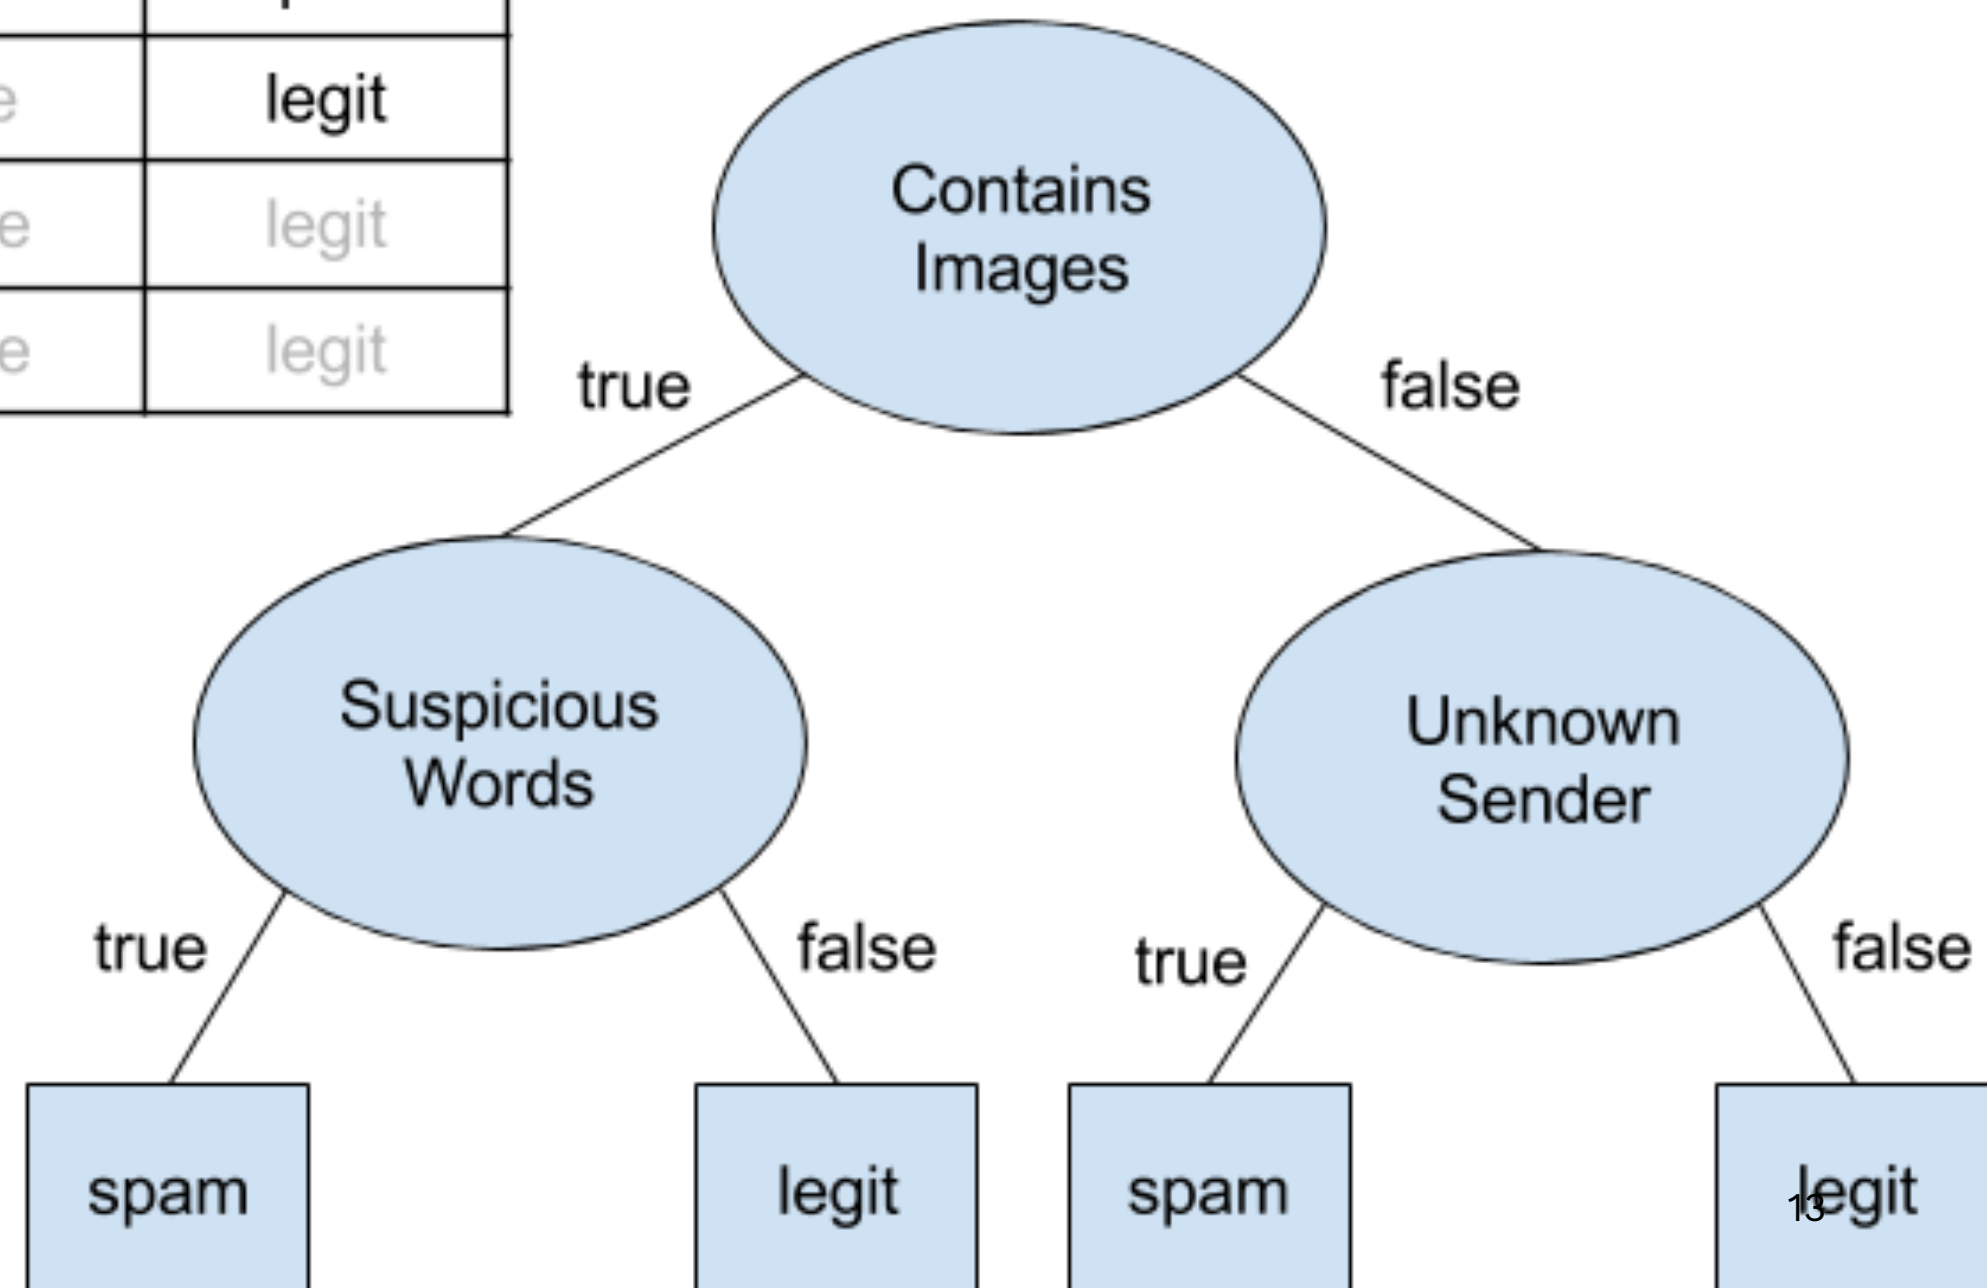
ID	Susp. words	Unknown sender	Contains images	Classification
1	true	false	true	spam
2	true	true	false	spam
3	true	true	true	spam
4	false	true	true	legit
5	false	false	false	legit
6	false	false	false	legit

Which remaining feature best discriminates the data?



ID	Susp. words	Unknown sender	Contains images	Classifi- cation
1	true	false	true	spam
2	true	true	false	spam
3	true	true	true	spam
4	false	true	true	legit
5	false	false	false	legit
6	false	false	false	legit

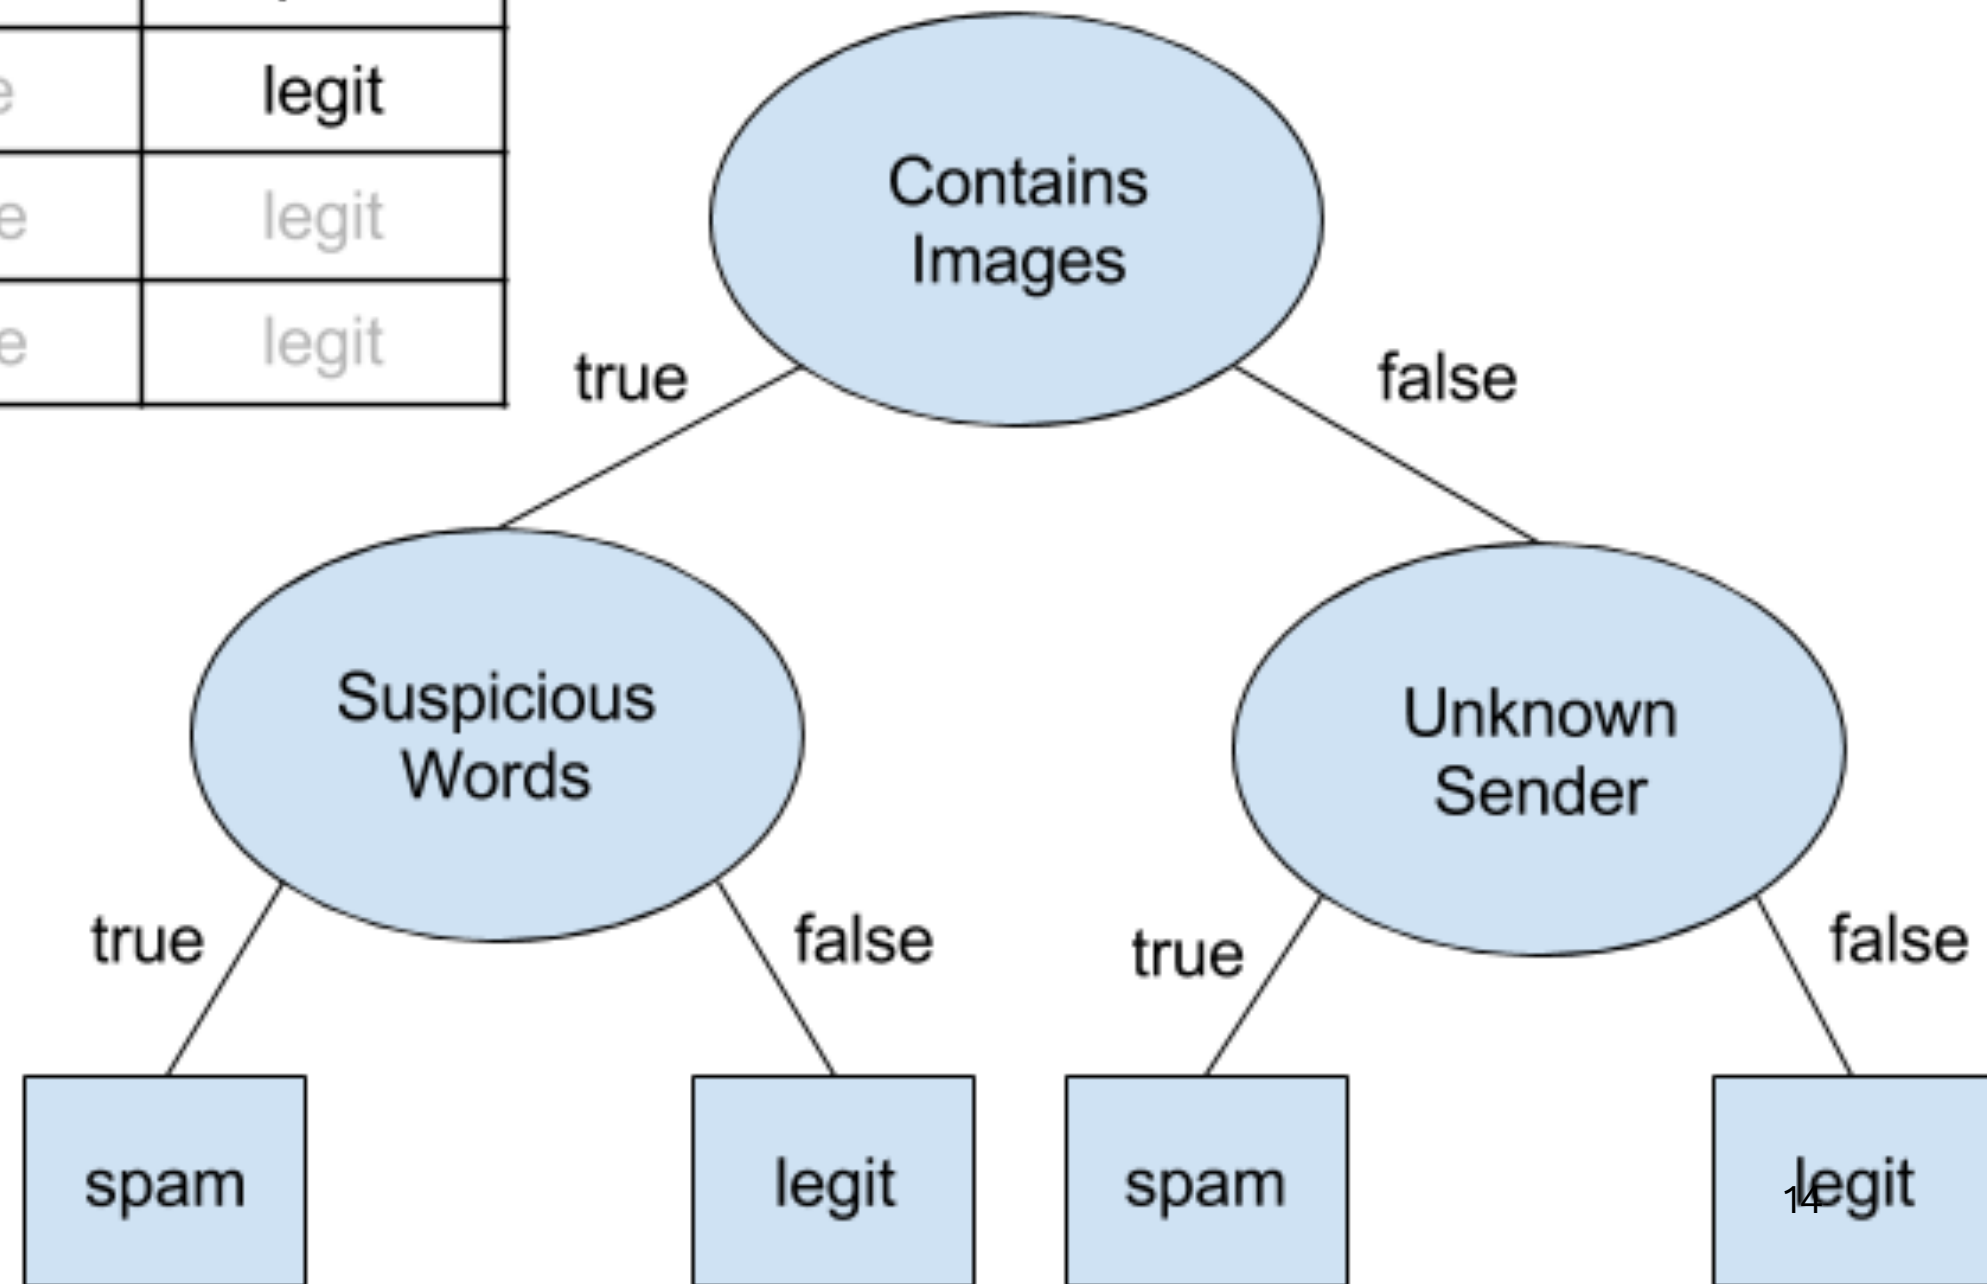
We can stop here!



ID	Susp. words	Unknown sender	Contains images	Classifi- cation
1	true	false	true	spam
2	true	true	false	spam
3	true	true	true	spam
4	false	true	true	legit
5	false	false	false	legit
6	false	false	false	legit

We can stop here!

*What would happen if we picked a different feature for the root node?*





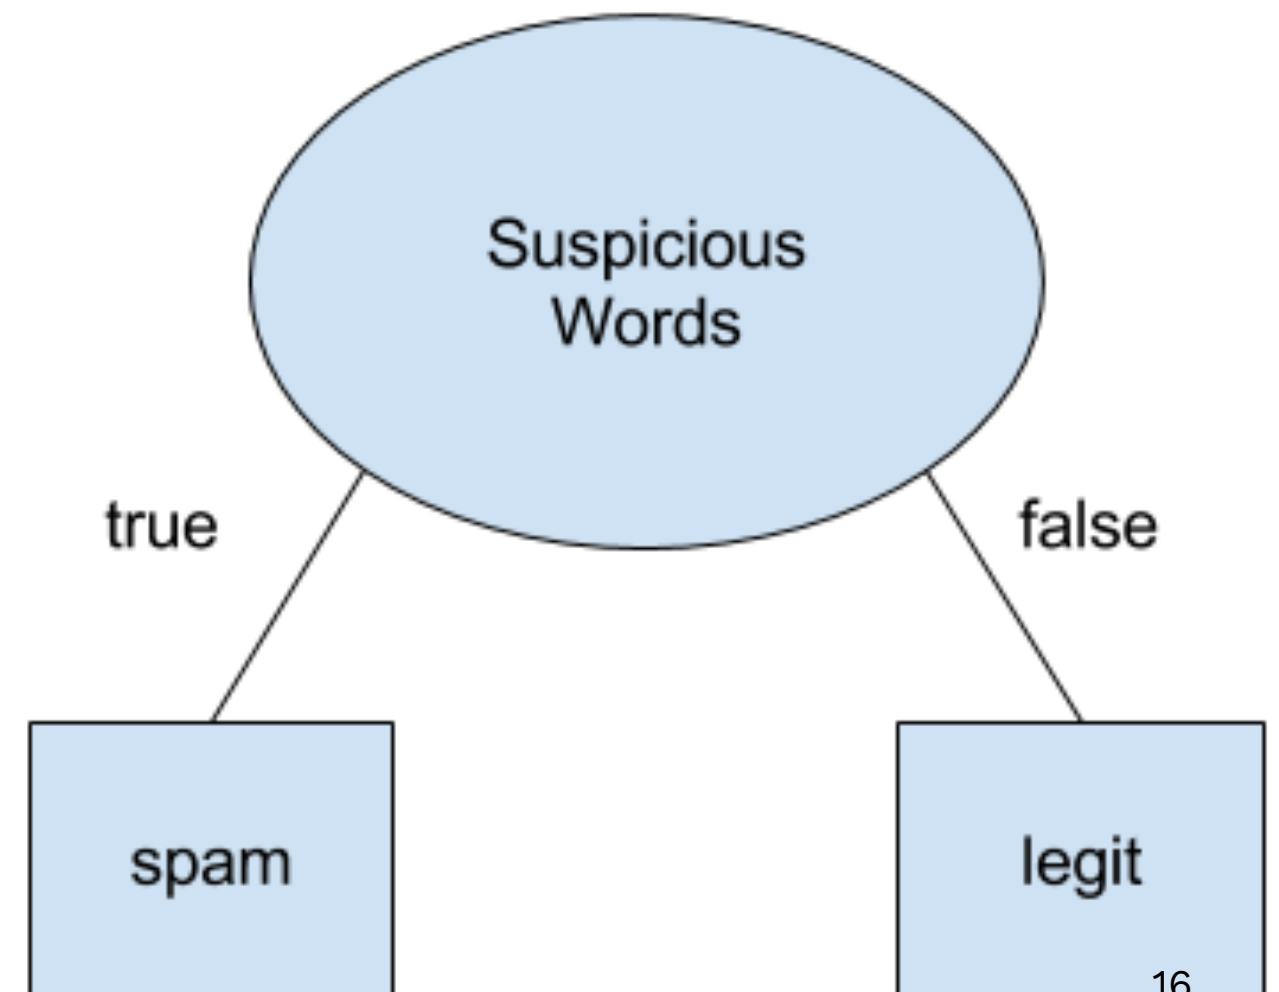
# Practice

Which feature *best* discriminates the data?

Draw the decision tree starting with this feature as the root.

<i>ID</i>	<i>Susp. words</i>	<i>Unknown sender</i>	<i>Contains images</i>	<i>Classifi- cation</i>
1	true	false	true	spam
2	true	true	false	spam
3	true	true	true	spam
4	false	true	true	legit
5	false	false	false	legit
6	false	false	false	legit

<i>ID</i>	<i>Susp. words</i>	<i>Unknown sender</i>	<i>Contains images</i>	<i>Classifi- cation</i>
1	true	false	true	spam
2	true	true	false	spam
3	true	true	true	spam
4	false	true	true	legit
5	false	false	false	legit
6	false	false	false	legit





# Building Decision Trees

- You've seen that decision trees **split** the data at each node
- For each split we can calculate **purity**

# Building Decision Trees

$$\begin{aligned}\text{purity of class } i &= p(\text{class } i | \text{data at node}) \\ &= p(i | D)\end{aligned}$$

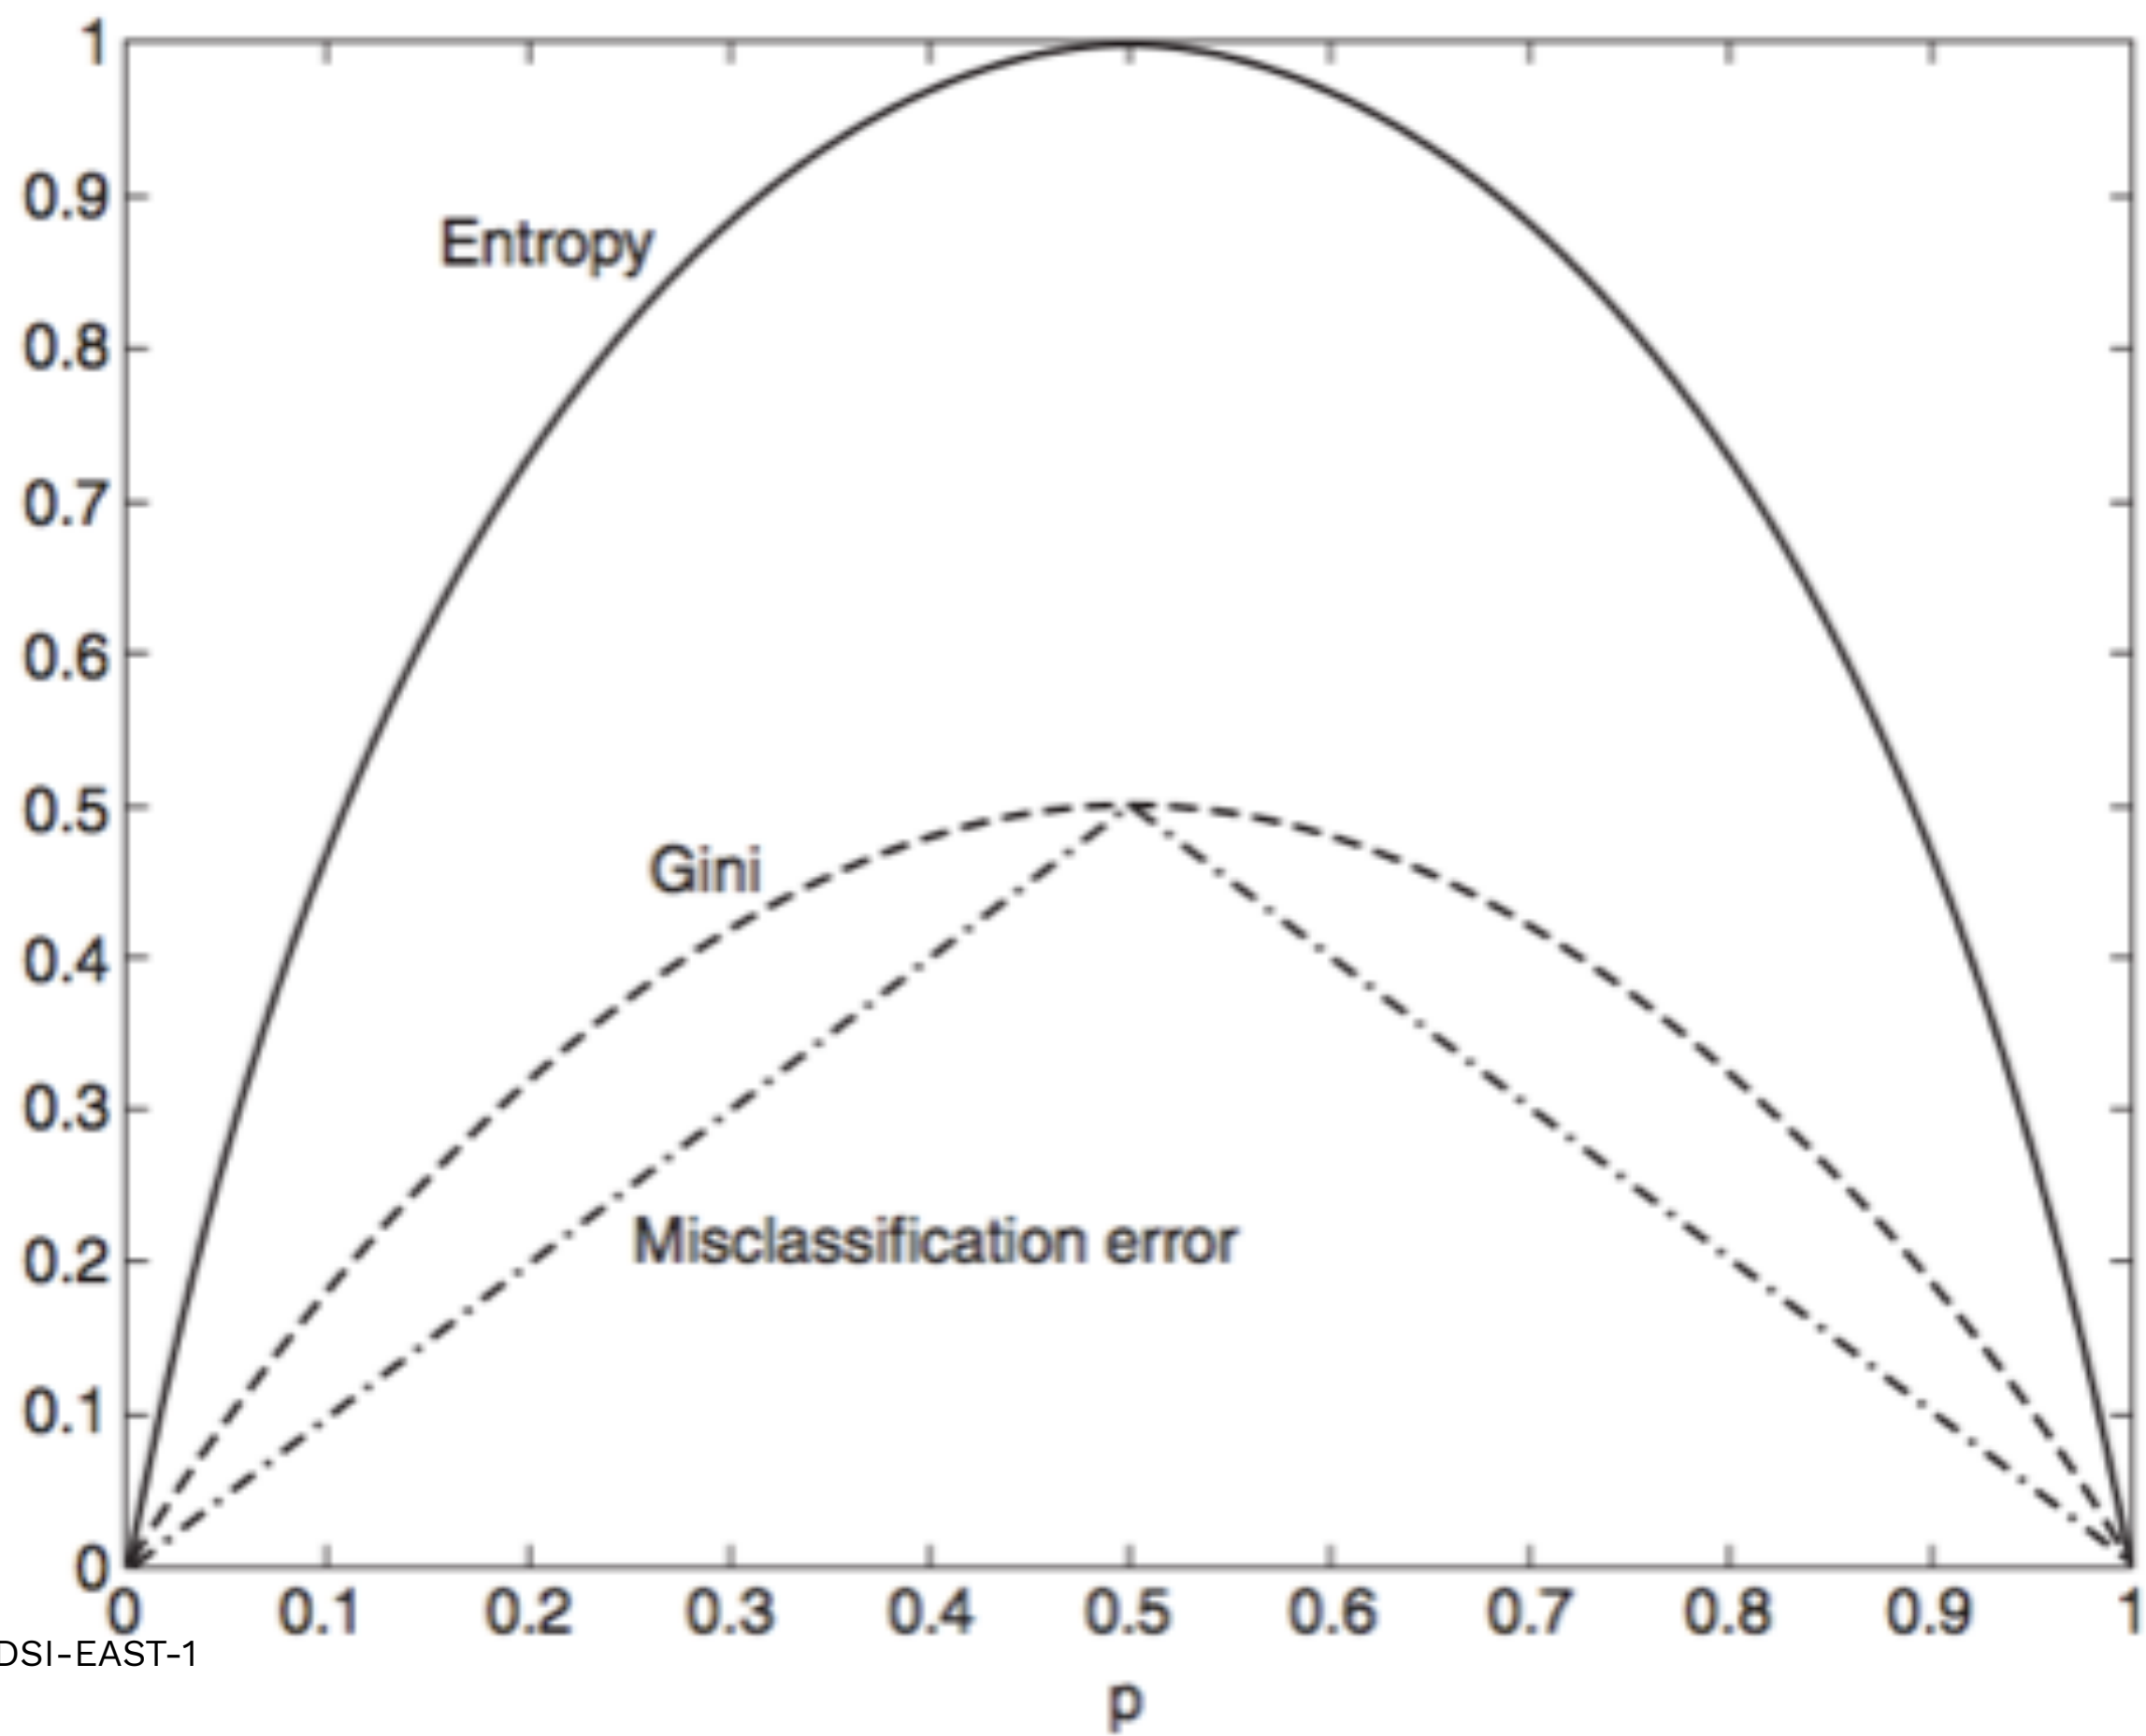
- For binary classification:
  - Worst case, purity = 0.5
  - Best case, purity = 1.0

# Generalizing Purity

Impurity measures for a node with data  $D$

$$\text{Entropy} = - \sum_{i=1}^{\text{classes}} p(i \mid D) \log_2 p(i \mid D)$$

$$\text{Gini} = \sum_{i=1}^{\text{classes}} p(i \mid D)(1 - p(i \mid D)) = 1 - \sum_{i=1}^{\text{classes}} p(i \mid D)^2$$



# Information Gain

**Goal:** determine how good a split is

**Solution:** gain

$$\text{gain} = I(\text{parent}) - \sum_{\text{children}} \frac{N_j}{N} I(\text{child}_j)$$

where  $I$  is the impurity measure,  $N_j$  is the number of observations at child node  $j$ , and  $N$  is the number of observations at the parent node.

# Full CARTs

- So far we've seen trees for **classification**
- Decision trees can be used for **regression** too

# Decision Trees

## Notes and Observations

- Decision trees are **hierachical**
  - Sequence of "if-this-than-that" conditions
- Decision trees are **non-parametric**
  - No  $\beta$  coefficients!
  - No assumption on distributions

# Decision Trees

## Notes and Observations

### CART advantages

- Simple to understand and interpret.
- Requires little data preparation.
- Able to handle both numerical and categorical data.
- Possible to validate a model using statistical tests.
- Once trained can be implemented on hardware and has extremely fast (real-time) execution.



# Decision Trees

## Notes and Observations

### CART disadvantages

- Locally-optimal.
- Overfitting.
- There are concepts that are hard to learn (XOR, parity or multiplexer problems)
- Decision trees can be biased if some classes dominate.